


Closed-form recovery of speech features from one federated update



Past the first-layer limit, deeper layers leak $\text{span}\{[h(c_t), 1]\}$ in 2,048 to 4,096 dimensions; frames are then decoded one at a time, each checked against the span.